

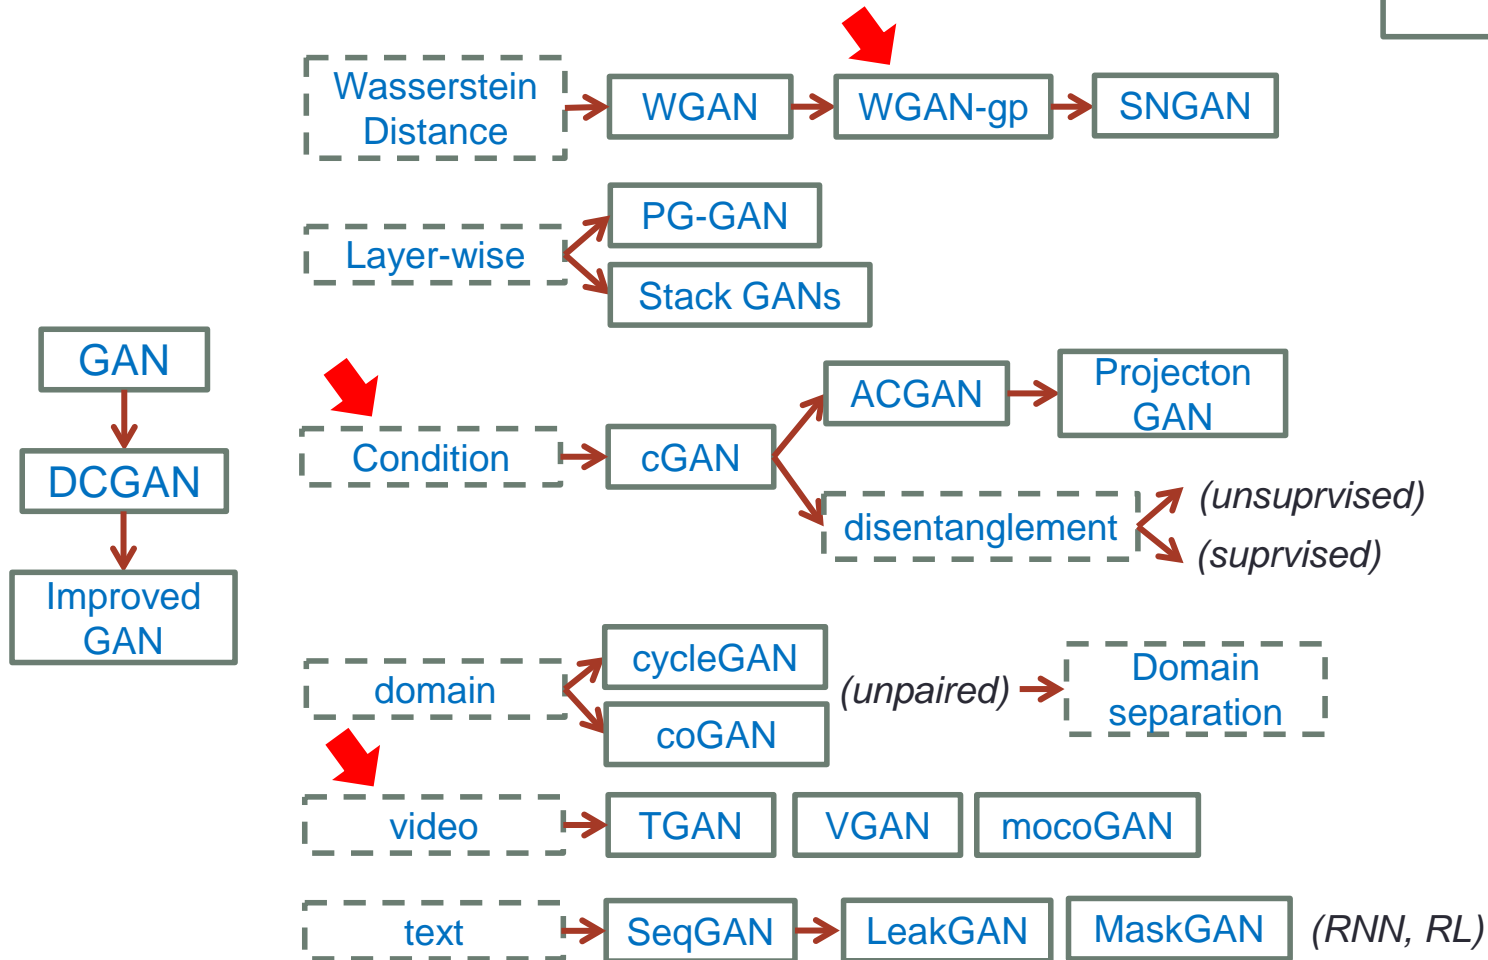
GAN

2018/7/11

Preface

- Topics
 - Experiences & tricks
 - Paper Survey
 - MuseGAN, MidiNet Review
- Outline
 - GAN Review
 - Road Map
 - Condition
 - Video

Road Map



GAN Review

Data & Generator

Discriminator

Training

Testing

Discussion

Review:

How to start a GAN project?

- Data
- Generator
- Discriminator
- Training
- Testing
- **State-of-the-art?**

↑ [-] **ajmooch** 6 points 3 months ago

↓ Projection discriminator + SN-GAN + progressive growing is probably our best bet at the moment for highest-fidelity high-res images, but the resources to do an ImageNet-level variety of classes at high res will probably be pretty substantial (just to get everything tuned, let alone to train it).

- **Models:** DCGAN WGAN-gp SNGAN

Data & Generator

- No matter what your task is, **Normalization** on data is necessary.
- The activation of the output layer of the generator depends on the range of the data

- Bounded

$[-1, 1]$ \rightarrow tanh
 $[0, 1]$ \rightarrow sigmoid


- Unbounded

zero means unit variance
logarithm (ex: on spectrogram)

leaky ReLU, ReLU?




- SharedArray & Shuffle data by index at runtime
- **z**: sampled from a **Gaussian distribution**



Generator

- **Batch Normalization** is essential to the quality
- Upsampling Layers:
 - Deconvolution (transposed convolution)
 - The most common method
 - Checkerboard artifacts?
 - Resize-Convolution
 - ex: PG-GAN uses nearest neighbor upscaling
 - Pixel Shuffling
 - Super resolution



Salimans et al., 2016 [2]

- Residual Blocks can enhance the quality as well, especially when the size of the images is large and the corresponding network is deep.

* **PG-GAN:** <https://arxiv.org/pdf/1710.10196.pdf>

* <https://distill.pub/2016/deconv-checkerboard>

Discriminator

DCGAN	WGAN-gp	SNGAN
<ul style="list-style-type: none">• BN applied• Sigmoid output	<ul style="list-style-type: none">• No BN• No Sigmoid	<ul style="list-style-type: none">• No BN• Spectral Norm• No Sigmoid

- Loss function:

Original GAN loss: DCGAN

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

D Loss: $-\log(D(\mathbf{x})) - \log(1 - D(G(\mathbf{z})))$

G Loss: $-\log(D(G(\mathbf{z})))$

Minimize!!



In practice...

Discriminator

WGAN loss: SNGAN

$$\text{D Loss: - } (D(X)) + (D(G(z)))$$

$$\text{G Loss: - } (D(G(z)))$$

Minimize!!

remove log

Gradient penalty: WGAN-gp

$$\text{D Loss: - } (D(X)) + (D(G(z))) + \frac{(\|\nabla_x D(x)\| - 1)^2}{\text{penalty}}$$

$$\text{G Loss: - } (D(G(z)))$$

penalty

Minimize!!














- Parameter Updating:


- WGAN-gp update **G** once every 5 updates of **D**

Training

- When to stop?
- How to monitor the training procedure?
- GAN Losses **cannot** truly reflect the quality (even for WGAN)
- Generate samples along training!

Samples/*.png:



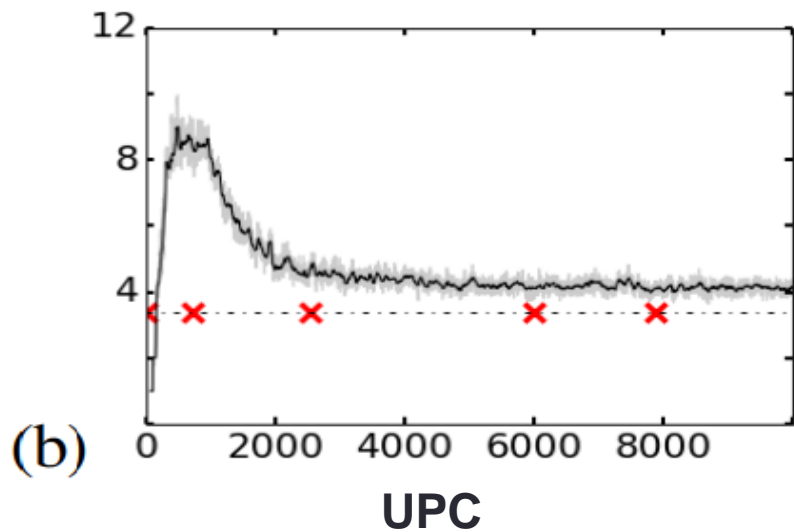
 sample_2.png	2018/7/7 下午 06:42
 sample_142.png	2018/7/7 下午 06:43
 sample_282.png	2018/7/7 下午 06:43
 sample_422.png	2018/7/7 下午 06:43
 sample_562.png	2018/7/7 下午 06:43
 sample_702.png	2018/7/7 下午 06:43
 sample_842.png	2018/7/7 下午 06:44
 sample_982.png	2018/7/7 下午 06:44
 sample_2002.png	2018/7/7 下午 06:45
 sample_4002.png	2018/7/7 下午 06:48
 sample_6002.png	2018/7/7 下午 06:50
 sample_8002.png	2018/7/7 下午 06:53
 sample_10002.png	2018/7/7 下午 06:55



(SNGAN on CIFAR10)

Early Stopping

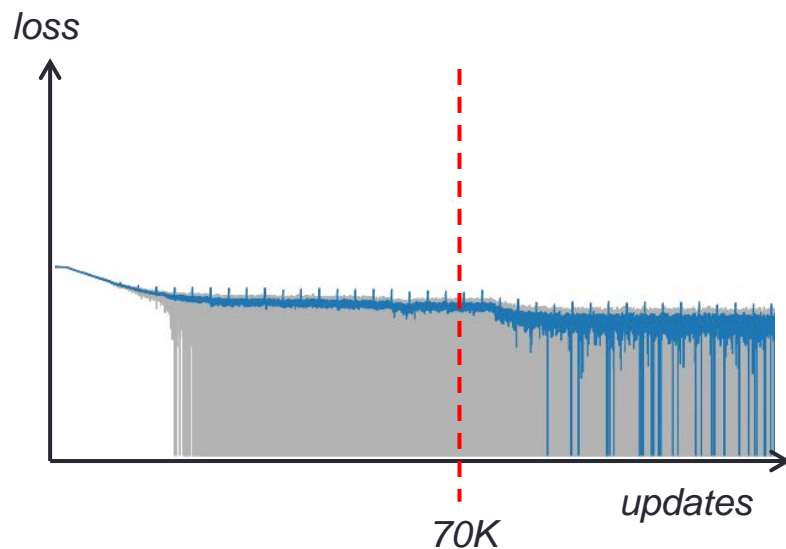
- Evaluate along training!




* from MuseGAN

Training


- For the first time, train the model for a longer period of time and monitor the procedure.
- Possible failure:
Mode collapse
Degeneration



(70K updates)





(130K updates)



Testing

- Interpolation: Spherical (instead of linear)




- Batch normalization **off**
'is_training' in TensorFlow
- Dropout **on?**
 - Image-to-Image (<https://arxiv.org/pdf/1611.07004v1.pdf>) turns on when testing
 - Later works seldom use it

* **GAN hacks:** <https://github.com/soumith/ganhacks>


* **Figure:** from https://github.com/ptrblck/prog_gans_pytorch_inference

Discussion

- So, which model should I use?
- Loss
 - For image, **SNGAN** outperforms others in both of the efficiency and the quality.
 - For discrete (binary) representation, only **WGAN-gp** can generate results successfully.
- Network design
 - Deconvolution is simple but powerful
 - If you aim to generate images with higher qualities, try to use residual blocks and resize-upsampling layers.



(WGAN-gp on Lead Sheet)



(SNGAN on Lead Sheet)


Conditional GANs

How to apply condition
Tag? Disentanglement
Discussion


How to apply condition?

- Input Concatenation
 - cGAN: <https://arxiv.org/pdf/1411.1784.pdf>
- Feature Map Concatenation
 - DCGAN: <https://github.com/carpedm20/DCGAN-tensorflow>
- Auxiliary Classifier
 - ACGAN: <https://arxiv.org/pdf/1610.09585.pdf>
- Encoder
 - S^2 GAN, Generative Adversarial Text to Image Synthesis
 - MuseGAN/MidiNet: <https://arxiv.org/pdf/1703.10847.pdf>
 - FTGAN: <https://arxiv.org/pdf/1711.09618.pdf>
- Projection Discriminator (ICLR, 2018)
 - <https://arxiv.org/pdf/1802.05637.pdf>
- Principle:
 - Both of G and D need to receive the conditional information.


**(a) cGANs,
input concat**
(Mirza & Osindero, 2014)



**(b) cGANs,
hidden concat**
(Reed et al., 2016)



(c) AC-GANs
(Odena et al., 2017)



(d) (ours) Projection




Figure 1: Discriminator models for conditional GANs

Concatenation



- The most common method, naive but practical.

vector concatenation (for linear layer, input z)



**one-hot
label**

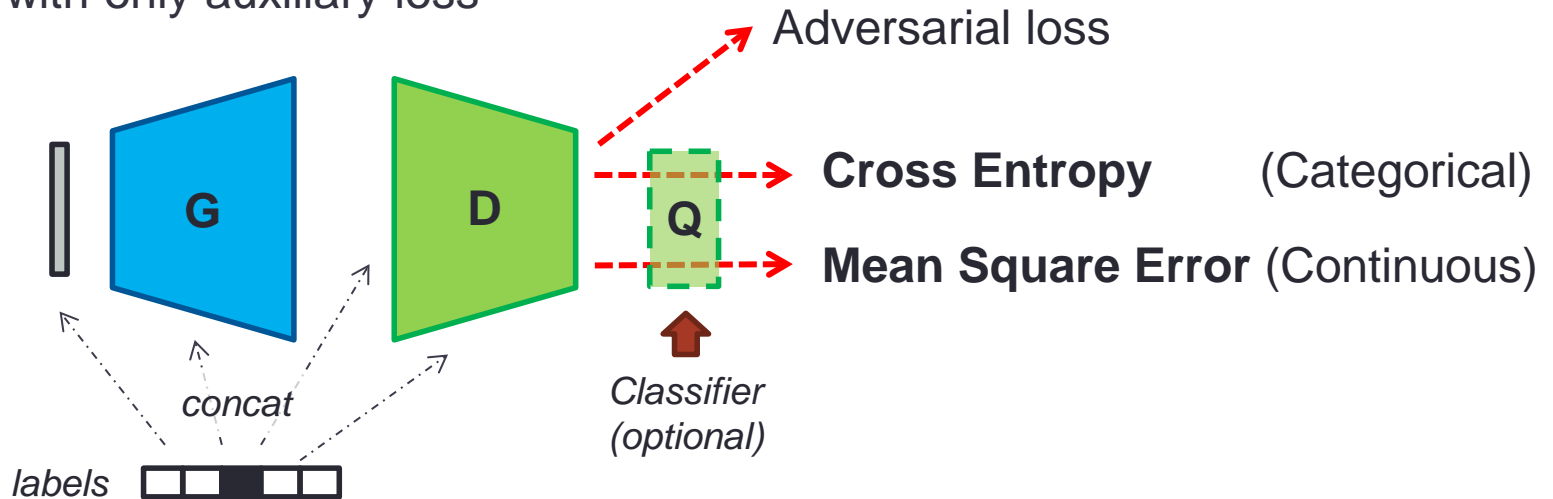
Feature map concatenation (for convolution layer)



(SNGAN on MNIST)

Auxiliary Classifier

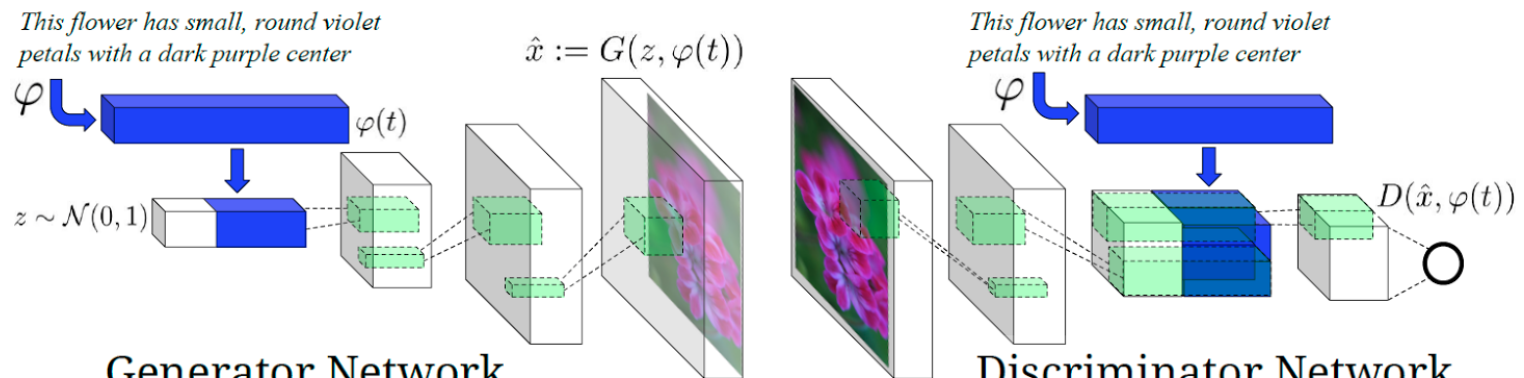
- Or with only auxiliary loss




- Where should I concatenate labels? Every layer in **G** & **D**?
 - It depends. Layers which is responsible for higher level features have priorities.
- Conditional generation is still a challenging task.
- SNGAN + projection discriminator achieve promising results.

Encoder

- For more complex conditional vectors
- Motivation: guide the generation process




Generative Adversarial Text to Image Synthesis



S² GAN: Generative Image Modeling using Style and Structure Adversarial Networks

U-Net Encoder

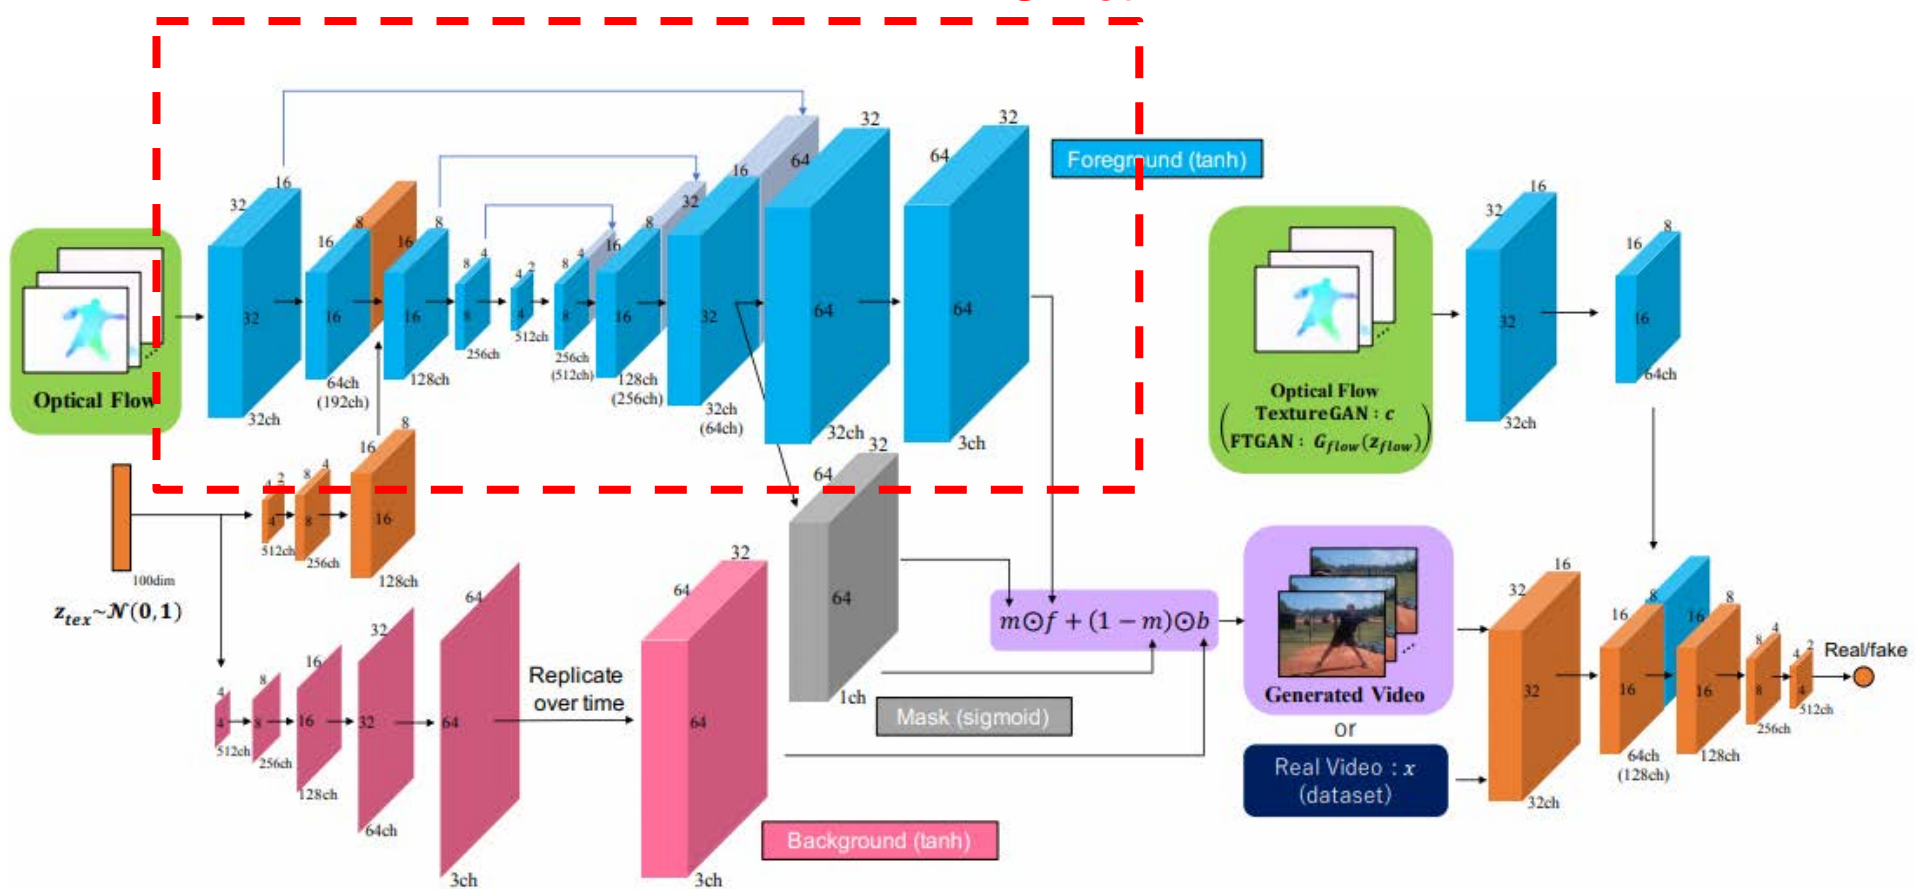
- Skip-connection: no information loss
- MidiNet: previous bar
- MuseGAN: accompanied track



U-Net Encoder

- FTGAN (video generation)
- Use conditional vector to guide the generation

U-Net



* *Optical flow: a feature about motion*

Tag

- **Supervised** GAN training
- **Conditional generation**
 - Given a class, generate images according to that
- **Disentanglement**
 - Try to acquire the **attribute-invariant** latent space or the controllability **with limited labels** (Ex: Gender)
 - Related Works:
 - **DR-GAN**
 - **TD-GAN**
 - **Fader Network**
 - **StarGAN** (CycleGAN)


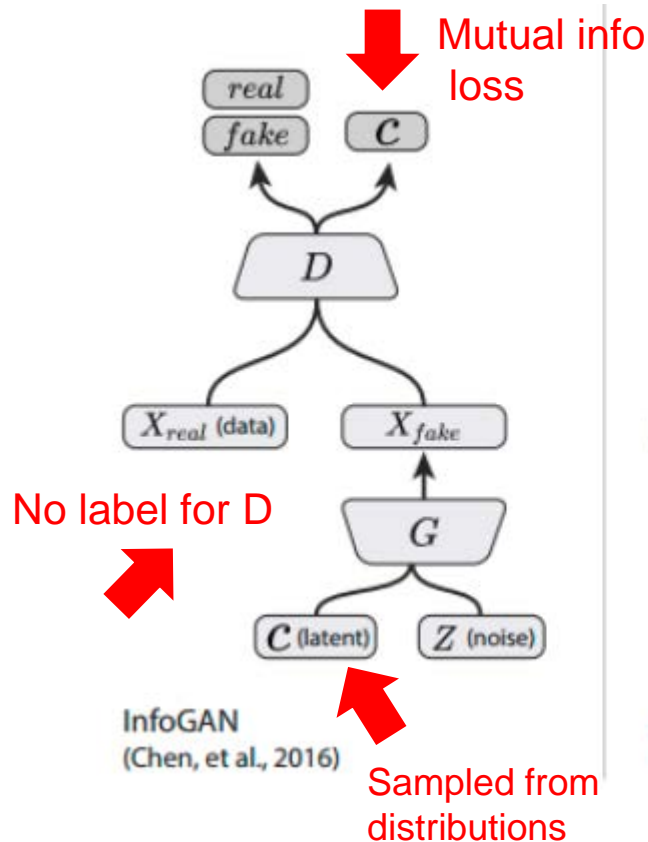


(male -> female) from Fader Network

Disentanglement

- Unsupervised?
 - infoGAN

add mutual information loss to encourage latent codes learn the most obvious properties



(a) Varying c_1 on InfoGAN (Digit type)
(categorical)



(c) Varying c_2 from -2 to 2 on InfoGAN (Rotation)
(continuous)

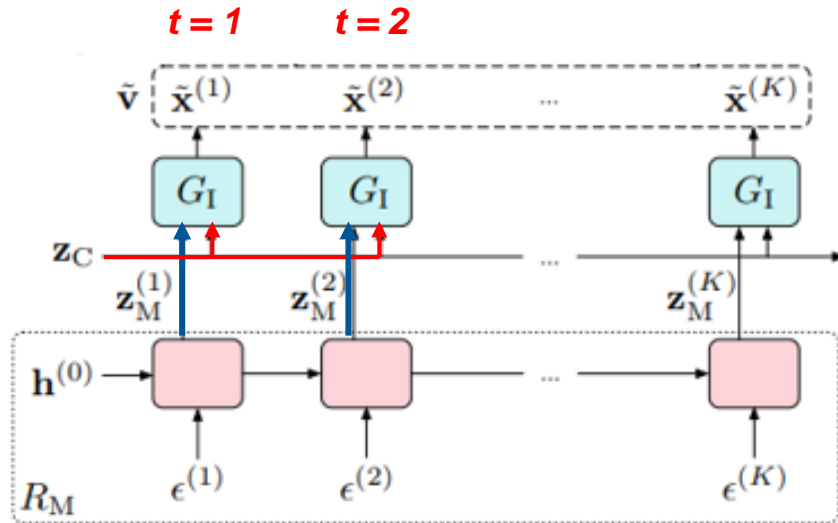
Disentanglement

- **Unsupervised**


- mocoGAN & MuseGAN

Design reasonable network architecture to encourage latent codes learn the **variant** and **invariant** properties

- Both of works simply use 3DCNN in **D**



mocoGAN



(c) Hybrid model

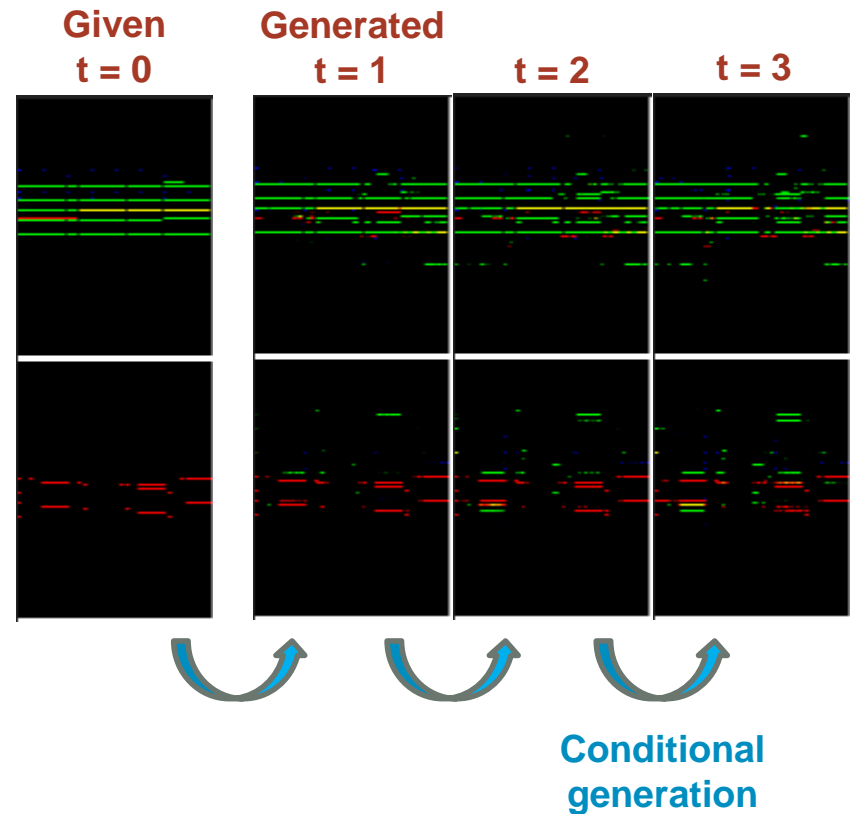
MuseGAN

Discussion

- Previous bar as condition (MidiNet)
 - For melody **Good**
 - For multi-track **Failed!**

- Reasons:
 - The source and target domain are the same
 - There are too many possible results

Finally, the network tries to copy and paste (AE, instead of U-Net)



Video GANs

Related Works
Discussion

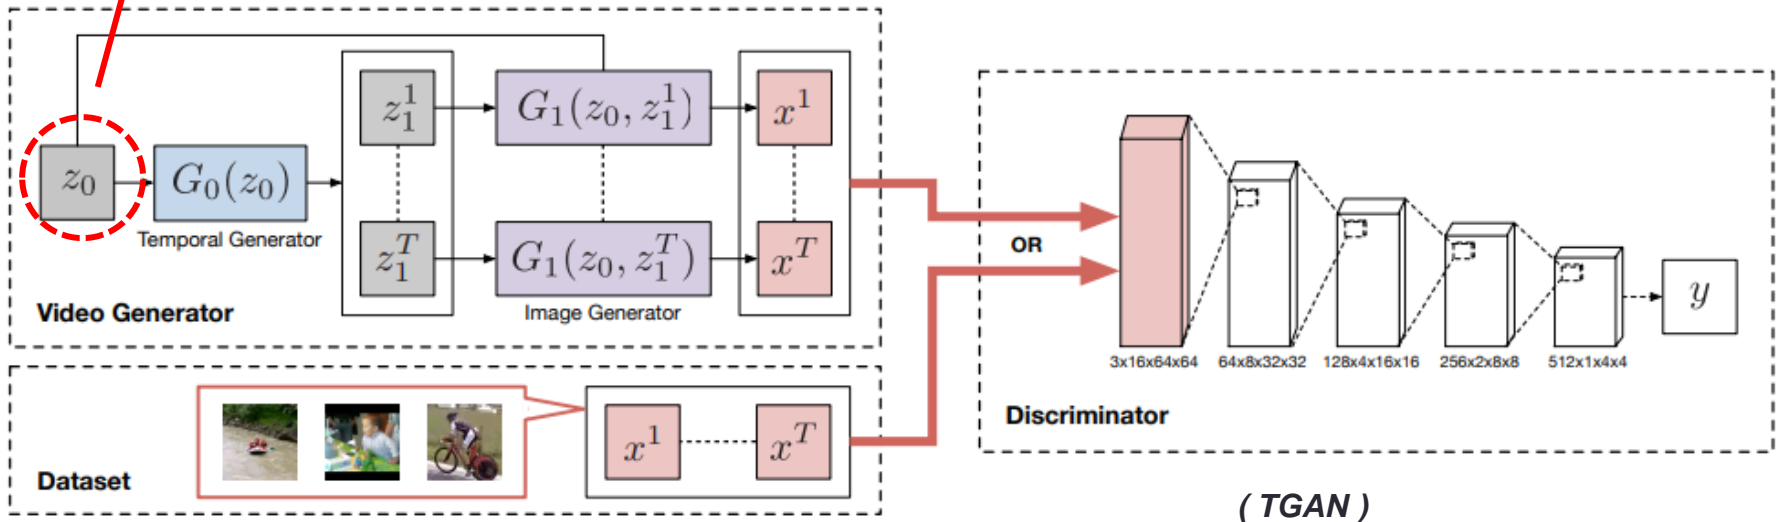

Video Generation

- Early papers try to predict the **next frame**
 - MidiNet
- Recent works
 - VGAN
 - TGAN
 - mocoGAN
 - FTGAN
- It's too difficult to generate a sequence of images directly.
 - Decomposition the video
 - VGAN: foreground and background
 - mocoGAN: content and motion
 - Using additional information to guide the generation
 - FTGAN: optical flow

Video Generation

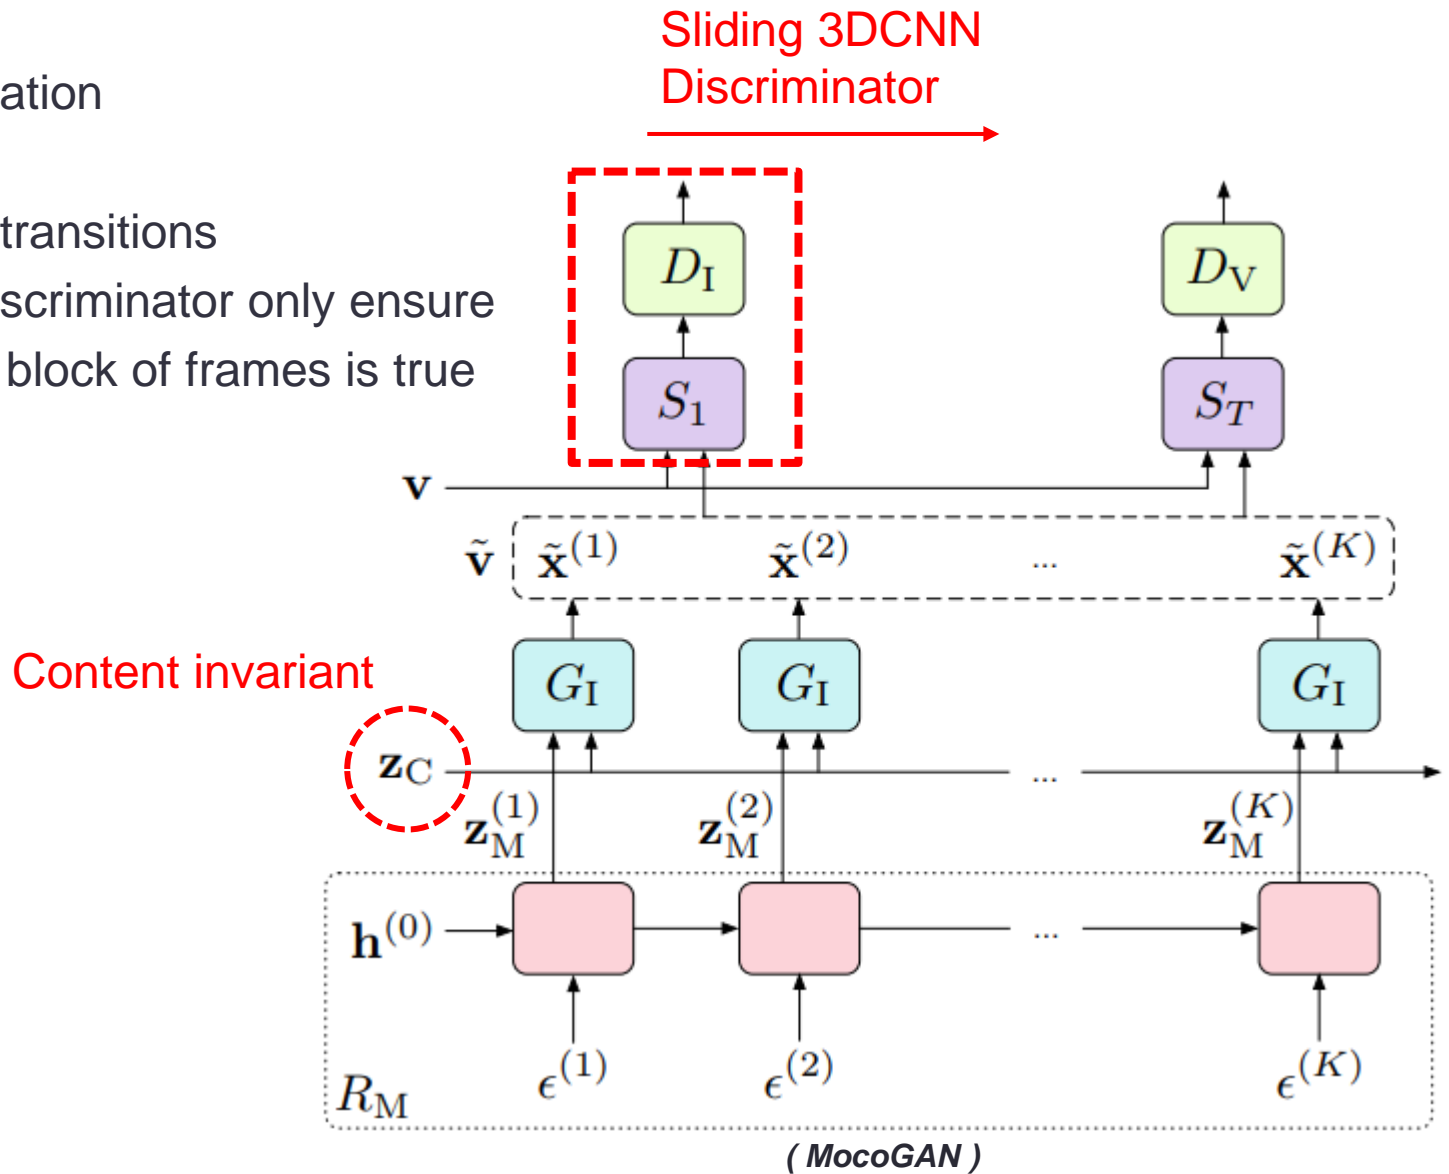
- Problems?

From one point at the latent space!



Video Generation

- GIF generation
- No scene transitions
- 3DCNN discriminator only ensure every local block of frames is true



conclusion

A thin, vertical red line is positioned to the right of the word "conclusion", extending from the top of the word down to the bottom of the slide.

Video Generation

- GAN templates:
 - Chainer:
 - <https://github.com/pfnet-research/chainer-gan-lib>
 - https://github.com/pfnet-research/sngan_projection
 - Tensorflow:
 - <https://github.com/carpedm20/DCGAN-tensorflow>
 - <https://github.com/wiseodd/generative-models>
 - Pytorch:
 - <https://github.com/eriklindernoren/PyTorch-GAN>

END

Text GANs

Related Works
Discussion